

Agentic Simulator

Extended Code Tutorial

A Discrete Event Simulator for LLM-Based Agentic Workloads
on Edge / Cloud Infrastructure

Thesis Project - Master's Degree

May 2026

Table of Contents

- 1. Overview & Motivation
- 2. Architecture at a Glance
- 3. Package Structure
- 4. Configuration Layer (config/AgenticConfig.java)
- 5. JSON Configuration File (agentic_config.json)
- 6. Event System (events/AgenticEvent.java)
- 7. Event Scheduler (scheduler/EventScheduler.java)
- 8. Stochastic LLM Engine (engine/LLMEngine.java)
- 9. Agent Decision Model (engine/AgentDecision.java)
- 10. Tool Pool (engine/ToolPool.java)
- 11. Agent Service (runtime/AgentService.java)
- 12. Orchestrator (runtime/Orchestrator.java)
- 13. Message Record (runtime/Message.java)
- 14. Trajectory Collector (metrics/TrajectoryCollector.java)
- 15. Simulation Runner (AgenticSimulationRunner.java)
- 16. Entry Point (SimulatorApplication.java)
- 17. End-to-End Flow Example
- 18. Key Design Decisions

1. Overview & Motivation

This project is a Discrete Event Simulator (DES) that models the behavior of LLM-based agentic systems running on edge/cloud infrastructure. Instead of calling real LLM APIs (which would be expensive and slow), it uses stochastic models to simulate realistic latency, cost, and decision-making behavior.

The simulator answers questions like:

- How does placing an agent on the edge vs. cloud affect end-to-end latency?
- What is the cost of running N concurrent agentic workflows per second?
- How do concurrency limits and queuing affect tail latency?
- What happens when tools fail or have high latency?

It is the v2 evolution of a Kubernetes microservice simulator (v1). While v1 modeled traditional request/response microservices with CPU time-slicing, v2 models agentic workflows where an LLM makes decisions in a loop: generate text, call tools, or delegate.

What makes it 'agentic'?

In a traditional microservice, the request follows a fixed path (Service A -> Service B -> C). In an agentic system, the LLM DECIDES the next step at each iteration. It might:

1. Generate a final text response (workflow completes)
2. Call an external tool (search, database, API) and loop back
3. Delegate to another agent
4. Fail

This creates a variable-length, non-deterministic execution path that the simulator models using weighted random decisions and stochastic latency sampling.

2. Architecture at a Glance

The simulator follows a classic Discrete Event Simulation pattern:

1. A priority queue (min-heap) holds future events sorted by time.
2. The event loop pops the earliest event and dispatches it.
3. Processing an event may schedule new future events.
4. Time is simulated (not real) - it jumps from event to event.

This is single-threaded and deterministic given the same random seed.

Component Wiring Diagram

```
AgenticSimulationRunner
|-- loads JSON config --> Topology
|-- creates LLMEngine (stochastic LLM surrogate)
|-- creates ToolPool (stochastic tool surrogate)
|-- creates EventScheduler (priority queue)
|-- creates TrajectoryCollector (CSV logger)
|-- creates Orchestrator (event router)
|   |-- registers AgentService(s)
|
|-- generates Poisson traffic (workflow submissions)
|-- runs event loop: scheduler.run(orchestrator::process)
|-- exports trajectory CSV
```

Event Flow

```
User Request (Poisson)
--> [SUBMIT] Orchestrator schedules AgentReceive
--> [AGENT_RECEIVE] AgentService builds context, calls LLM
--> schedules LlmComplete (after network + inference latency)
--> [LLM_COMPLETE] Branch on decision:
    GENERATE_TEXT --> workflow SUCCESS, terminate
    CALL_TOOL      --> schedule ToolComplete
    DELEGATE       --> (reserved for multi-agent)
    FAIL           --> workflow FAIL, terminate
--> [TOOL_COMPLETE] Result feeds back as new AgentReceive
--> loop continues until text/fail/budget exhausted
```

3. Package Structure

```
com.thesis.simulator.agentic/
|-- AgenticSimulationRunner.java    (main runner, wiring, traffic gen)
|
|-- config/
|   +-- AgenticConfig.java          (all data records: Topology, Zone,
|                                   LLMPProfile, ToolProfile, etc.)
|-- engine/
|   +-- LLMEngine.java              (stochastic LLM: latency, cost, decisions)
|   +-- AgentDecision.java          (decision enum + record)
|   +-- ToolPool.java               (stochastic tool calls)
|
|-- events/
|   +-- AgenticEvent.java           (sealed event hierarchy)
|
|-- runtime/
|   +-- Orchestrator.java           (event dispatcher + workflow submission)
|   +-- AgentService.java           (the agent loop: receive->LLM->branch)
|   +-- Message.java                (token-carrying message record)
|
|-- scheduler/
|   +-- EventScheduler.java         (priority queue event loop)
|
+-- metrics/
    +-- TrajectoryCollector.java     (in-memory log, CSV export)
```

The separation is intentional:

- 'config' holds pure data (records) with no behavior.
- 'engine' holds the stochastic models (LLM, tools) - the 'math'.
- 'events' defines the message types flowing through the system.
- 'runtime' holds the processing logic (agent loop, orchestrator).
- 'scheduler' is the DES engine itself (reusable, generic).
- 'metrics' handles observability and output.

4. Configuration Layer

File: config/AgenticConfig.java

This file contains only Java records (immutable data classes). No behavior, just structure. It defines the entire topology of a simulation experiment.

Zone

```
public record Zone(String id) {}
```

A logical placement zone like "CLOUD", "EDGE", or "FOG". Agents, LLMs, and tools are each hosted in a zone. Network latency depends on zone pairs.

NetworkLink

```
public record NetworkLink(
    String fromZone, String toZone,
    double latencyMs, double bandwidthMbps
) {}
```

A directed network link. When an agent in zone A calls an LLM in zone B, the simulator looks up the (A->B) link to add network latency. Every cross-zone call pays round-trip: A->B + B->A.

LLMProfile

```
public record LLMProfile(
    String id,
    double ttftMeanMs,           // Time To First Token (mean)
    double ttftGammaShape,      // Gamma distribution shape param
    double tpotMeanMs,          // Time Per Output Token (mean)
    double tpotGammaShape,      // Gamma distribution shape param
    int outputTokensMean,       // How many tokens per response
    int outputTokensStd,
    double pricePerInputToken,
    double pricePerOutputToken,
    String hostZone
) {}
```

Stochastic profile of an LLM model. The key insight: real LLM inference latency follows a Gamma distribution (always positive, right-skewed with occasional outliers). Total latency = TTFT + outputTokens * TPOT.

TTFT (Time To First Token) is the startup delay before streaming begins. TPOT (Time Per Output Token) is the per-token generation speed. Both are sampled from Gamma(shape, scale=mean/shape).

ToolProfile

```
public record ToolProfile(
    String id,
    double latencyMeanMs, double latencyStdMs,
```

```
    int responseTokensMean,
    double errorRate,
    String hostZone
) {}
```

Profile for external tools (search APIs, databases, MCP functions). Latency is Gaussian. Tools can fail with a configurable error rate.

AgentDefinition

```
public record AgentDefinition(
    String id,
    String llmProfileId,      // which LLM this agent uses
    List<String> toolIds,     // which tools are available
    String hostZone,
    int maxConcurrency,      // thread pool size
    double coldStartMs       // container startup penalty
) {}
```

Defines an agent service. Each agent is hosted in a zone, uses a specific LLM, has access to certain tools, and has a concurrency limit (like a Pod's thread pool). When all slots are busy, new requests queue up.

Topology & WorkflowSpec

```
public record Topology(
    List<Zone> zones,
    List<NetworkLink> links,
    Map<String, LLMPProfile> llmProfiles,
    Map<String, ToolProfile> tools,
    Map<String, AgentDefinition> agents,
    WorkflowSpec workflow
) {}

public record WorkflowSpec(
    String entryAgent, int initialPromptTokens,
    int maxSteps, int maxTokens
) {}
```

Topology bundles the entire experiment configuration. WorkflowSpec defines safety limits: max steps per workflow (prevents infinite loops) and max tokens (budget cap).

5. JSON Configuration File

File: src/main/resources/agentic_config.json

This JSON drives the entire simulation. Here is the current config with explanations:

Nodes & Zones

```
"nodes": [{
  "nodeId": 1, "zone": "CLOUD",
  "cores": 8, "frequencyHz": 3e9,
  "vramGB": 80, "bandwidthBytesPerSec": 1.25e9
}]
```

Defines physical/virtual nodes. Each node belongs to a zone (CLOUD, EDGE, etc.).

Zone Pair Latencies

```
"zonePairLatencies": [
  { "from": "CLOUD", "to": "CLOUD", "latencyMs": 1.0 }
]
```

Network latency between zones. CLOUD->CLOUD is 1ms (intra-datacenter). For multi-zone experiments, you would add entries like EDGE->CLOUD: 50ms, CLOUD->EDGE: 50ms.

LLM Models

```
"llmModels": [{
  "modelId": "gpt-4o",
  "ttftMeanMs": 400.0, "ttftGammaShape": 4.0,
  "tpotMeanMs": 25.0, "tpotGammaShape": 10.0,
  "outputTokensMean": 150.0, "outputTokensStdDev": 50.0,
  "costPerInputToken": 0.0000025,
  "costPerOutputToken": 0.00001,
  "hostZone": "CLOUD",
  "decisionWeights": { "generate_text": 0.4, "call_tool": 0.6 }
}]
```

This defines a GPT-4o-like model with:

- TTFT ~ Gamma(shape=4, mean=400ms) -> typical range: 100-800ms
- TPOT ~ Gamma(shape=10, mean=25ms) -> typical range: 15-40ms per token
- Output length ~ Gaussian(mean=150, std=50) tokens
- Cost: \$2.50/M input tokens, \$10/M output tokens
- Decision weights: 60% chance to call a tool, 40% to generate final text

Tools

```
"tools": [{
  "toolId": "search-tool",
  "latencyMeanMs": 200.0, "latencyStdDevMs": 50.0,
  "resultTokensMean": 300,
```



```
"errorRate": 0.0,  
"hostZone": "CLOUD"  
}]
```

A search tool with ~200ms latency returning ~300 tokens. 0% error rate.

Simulation Parameters

```
"simulation": {  
  "durationSeconds": 100.0,  
  "arrivalRate": 0.5,  
  "userMessageTokensMean": 50.0,  
  "userMessageTokensStdDev": 20.0,  
  "maxStepsPerWorkflow": 10,  
  "maxTokensPerWorkflow": 10000,  
  "seed": 42  
}
```

Run for 100 seconds of simulated time. Users arrive at 0.5 requests/second (Poisson process). Each user message is ~50 tokens. Workflows are capped at 10 agent steps or 10,000 tokens. Seed=42 for reproducibility.

6. Event System

File: events/AgenticEvent.java

This uses Java's sealed interface + records for type-safe, exhaustive event types. Every event has a `timeMs()` and is `Comparable` (sorted by time in the priority queue).

Event Types

TaskSubmit

Reserved placeholder for queued submissions. Not used in the current flow.

AgentReceive

```
record AgentReceive(double timeMs, String workflowId,
                   String agentId, Message message)
```

A message has arrived at an agent and is ready for processing. This is the entry point for every agent step. Carries the input tokens via `Message`.

LlmComplete

```
record LlmComplete(double timeMs, String workflowId,
                  String agentId, AgentDecision decision,
                  double costAccrued, int outputTokens)
```

The LLM has finished generating. The decision field tells the agent what to do next: generate text (finish), call a tool (loop), delegate, or fail.

ToolComplete

```
record ToolComplete(double timeMs, String workflowId,
                   String agentId, String toolId,
                   int responseTokens, boolean errored)
```

A tool call has returned. If `errored=true`, the workflow terminates. Otherwise, the response tokens feed back as a new `AgentReceive`.

WorkflowComplete

```
record WorkflowComplete(double timeMs, String workflowId,
                      String reason, double totalLatencyMs,
                      double totalCostUsd, int totalSteps)
```

Terminal event. Reason can be: `SUCCESS`, `BUDGET_EXHAUSTED`, `TOOL_FAILURE`, or `LLM_FAIL`.

7. Event Scheduler

File: scheduler/EventScheduler.java

The heart of any DES. This is intentionally minimal - just a priority queue wrapper.

```
public class EventScheduler {
    private final PriorityQueue<AgenticEvent> queue = new PriorityQueue<>();
    private double now = 0.0;

    public void schedule(AgenticEvent ev) {
        // Cannot schedule events in the past
        if (ev.timeMs() < now) throw ...;
        queue.offer(ev);
    }

    public void run(Consumer<AgenticEvent> handler) {
        while (!queue.isEmpty()) {
            AgenticEvent ev = queue.poll(); // earliest event
            now = ev.timeMs();              // advance clock
            handler.accept(ev);              // process it
        }
    }
}
```

Key points:

- PriorityQueue sorts events by timeMs() (via Comparable).
- The run() loop is the main simulation loop - it processes ALL events.
- Processing one event (e.g., AgentReceive) typically schedules new future events (e.g., LlmComplete at time now + latency). This is how the simulation progresses.
- Single-threaded: no real concurrency, just simulated time.

8. Stochastic LLM Engine

File: engine/LLMEngine.java

This is the mathematical core. It does NOT call any real LLM API. Instead, it samples from probability distributions that match real-world LLM behavior.

Latency Model

```
latency = TTFT + outputTokens * TPOT

TTFT ~ Gamma(shape, scale = mean/shape)
TPOT ~ Gamma(shape, scale = mean/shape)
```

Why Gamma distribution?

- Always positive (latency cannot be negative)
- Right-skewed (occasional high-latency outliers, matching real LLM behavior)
- Two-parameter: shape controls the 'spread', mean sets the center
- When shape is large (e.g., 10), Gamma approaches Gaussian
- When shape is small (e.g., 2), there's a long right tail

The implementation uses Marsaglia & Tsang's (2000) rejection sampling method, which is the standard efficient algorithm for Gamma sampling.

Decision Model

```
public AgentDecision decide(int inputTokens, List<String> availableTools) {
    return decisionPolicy.apply(inputTokens, availableTools);
}
```

The decision policy is a pluggable BiFunction injected via the constructor. The current implementation is a weighted random chooser:

- Roll a random number [0, 1)
- Walk through the cumulative weights (e.g., call_tool=0.6, generate_text=0.4)
- Return the first action whose cumulative weight exceeds the roll

This means with the default config: 60% of LLM steps call a tool, 40% finish.

Cost Model

```
cost = inputTokens * pricePerInputToken + outputTokens * pricePerOutputToken
```

Straightforward token-based pricing. Accumulated per workflow.

9. Agent Decision Model

File: engine/AgentDecision.java

```
public record AgentDecision(Kind kind, String targetId, int outputTokens) {
    public enum Kind {
        GENERATE_TEXT, // Final answer -> workflow completes
        CALL_TOOL,      // Invoke a tool -> loop back
        DELEGATE,        // Hand off to another agent
        FAIL             // LLM failure
    }
}
```

Simple value object representing what the LLM decided to do:

- GENERATE_TEXT: The agent produces a final answer. targetId is null.
- CALL_TOOL: targetId = the tool ID (e.g., 'search-tool').
- DELEGATE: targetId = another agent's ID. (Not yet implemented in v2.)
- FAIL: Simulates LLM failure/hallucination.

Factory methods (text(), tool(), delegate(), fail()) create instances cleanly.

10. Tool Pool

File: engine/ToolPool.java

Stochastic surrogate for external tool calls. Much simpler than LLMEngine.

```
sampleLatencyMs(toolId):
    latency ~ Gaussian(mean, stdDev), clamped >= 0

sampleError(toolId):
    return random() < errorRate

sampleResponseTokens(toolId):
    return profile.responseTokensMean (fixed, not random)
```

Tool latency uses Gaussian (can go near zero, unlike LLM which uses Gamma). Error sampling is a simple Bernoulli trial. Response tokens are currently fixed (could be made stochastic in a future version).

11. Agent Service - The Agent Loop

File: runtime/AgentService.java

This is the most complex class - it implements the full agentic processing loop. Think of it as a simulated microservice container running an LLM agent.

Concurrency Model

```
private int activeRequests = 0;
private final Queue<AgenticEvent.AgentReceive> waitingQueue = new LinkedList<>();
```

Each agent has maxConcurrency slots (like a thread pool). When all slots are busy, new requests enter a FIFO queue. When a workflow completes, it frees a slot and the next queued request is dequeued. This mirrors real container behavior (Tomcat/Jetty thread pools in Kubernetes Pods).

Per-Workflow State: WorkflowContext

```
private static class WorkflowContext {
    final String workflowId;
    final double startedAtMs;
    int stepIndex = 0;
    int accumulatedInputTokens = 0;
    int accumulatedOutputTokens = 0;
    double totalCostUsd = 0.0;
}
```

Each workflow gets its own context tracking: step count, accumulated tokens, and cost. This is keyed by workflowId in a HashMap. Context is created on first receive and destroyed on termination.

Step-by-Step Processing

Step 1: onReceive(AgentReceive)

Entry point. Checks concurrency:

- If at capacity: queue the request, log QUEUE_ENTER, return.
- Otherwise: increment activeRequests, proceed to processReceive().

Step 2: processReceive()

- Get or create WorkflowContext for this workflowId.
- Increment step counter, accumulate input tokens.
- Budget check: if steps > maxSteps or tokens > maxTokens, terminate.
- Look up the LLM profile (from agent's llmProfileId).
- Calculate network latency: agent zone -> LLM zone -> agent zone.
- Sample output tokens, decision, inference latency, and cost from LLMEngine.
- Schedule LlmComplete event at time: now + networkOut + inference + networkBack.

Step 3: onLlmComplete(LlmComplete)

The LLM has responded. Branch on the decision:

- GENERATE_TEXT -> terminate with SUCCESS
- CALL_TOOL -> dispatchTool()
- DELEGATE -> terminate with DELEGATE_NOT_IMPLEMENTED
- FAIL -> terminate with LLM_FAIL

Step 4: dispatchTool()

- Look up tool profile for network zone.
- Calculate network latency: agent zone -> tool zone -> agent zone.
- Sample tool latency and error flag.
- Schedule ToolComplete at time: now + networkOut + toolMs + networkBack.

Step 5: onToolComplete(ToolComplete)

- If errored: terminate with TOOL_FAILURE.
- Otherwise: create a follow-up Message with the tool's response tokens.
- Call processReceive() directly (no re-check concurrency - this workflow already holds a slot). This creates the LOOP back to Step 2.

Termination: terminate()

- Schedule WorkflowComplete event with total latency, cost, and step count.
- Remove WorkflowContext.
- Decrement activeRequests (free concurrency slot).
- Call tryDequeue() to process waiting requests.

12. Orchestrator

File: runtime/Orchestrator.java

Simple event router and workflow entry point. Two responsibilities:

1. submit(workflowId, promptTokens, now)

Creates the initial Message (from 'USER' to the entry agent) and schedules the first AgentReceive event. This is how workflows enter the system.

2. process(AgenticEvent)

Central dispatch method wired into the scheduler's main loop. Uses Java's pattern matching (instanceof) to route each event type to the correct AgentService method:

```
AgentReceive  -> agents.get(agentId).onReceive()
LlmComplete   -> agents.get(agentId).onLlmComplete()
ToolComplete  -> agents.get(agentId).onToolComplete()
WorkflowComplete -> log final metrics to TrajectoryCollector
```

13. Message Record

File: runtime/Message.java

```
public record Message(
    long id,           // auto-incrementing
    String fromAgent,  // sender ('USER', 'search-tool', etc.)
    String toAgent,    // receiver (agent ID)
    int inputTokens,   // token count (not bytes)
    Object payload,    // type-erased content
    double createdAtMs // timestamp
) {}
```

The data unit exchanged between entities. Key difference from v1's ServiceCall: it carries TOKEN counts, not byte counts. This is the fundamental unit of agentic workloads - everything is measured in tokens.

14. Trajectory Collector

File: metrics/TrajectoryCollector.java

In-memory event log that writes to CSV. Every significant event in the simulation is recorded with a flexible key-value payload.

```
void log(workflowId, entityId, eventType, timeMs, Map<String, Object> payload)

// Creates a Row(workflowId, entityId, eventType, timeMs, payload)
// Payload keys are dynamic (different per event type)
```

The writeCsv() method discovers ALL payload keys across ALL rows and writes them as columns. This means the CSV schema is self-adapting: adding new payload fields in the code automatically creates new CSV columns.

Example output columns:

time_ms, workflow_id, entity_id, event_type, step, input_tokens, decision, cost_usd, network_out_ms, inference_ms, ...

15. Simulation Runner

File: AgenticSimulationRunner.java

The main orchestration class. It wires everything together and runs the simulation. Here is what `run(configPath)` does, step by step:

Phase 1: Load Configuration

Reads the JSON config file using Jackson. Parses it into a `Topology` object containing all zones, links, LLM profiles, tool profiles, and agent definitions.

Phase 2: Build Decision Policy

Creates the weighted random decision function from the config's `decisionWeights`. This is a `lambda/BiFunction` that, given `(inputTokens, availableTools)`, rolls a random number and picks `call_tool`, `generate_text`, `delegate`, or `fail` based on cumulative probabilities.

Phase 3: Create Components

```
LLMEngine llm = new LLMEngine(new Random(seed), decisionPolicy);
ToolPool tools = new ToolPool(topology.tools(), new Random(seed+1));
EventScheduler scheduler = new EventScheduler();
TrajectoryCollector trace = new TrajectoryCollector();
Orchestrator orchestrator = new Orchestrator(topology, scheduler, trace);

// Register all agents
for each agent in topology:
    AgentService agent = new AgentService(def, topology, llm, tools, scheduler, trace);
    orchestrator.registerAgent(agentId, agent);
```

Phase 4: Generate Traffic (Poisson Process)

```
while (currentArrivalMs <= totalDuration){
    promptTokens = Gaussian(mean=50, std=20);
    orchestrator.submit('wf-NNNN', promptTokens, currentArrivalMs);

    // Exponential inter-arrival time
    gap = -ln(U) / lambda; // U ~ Uniform(0,1)
    currentArrivalMs += gap;
}
```

Poisson process: events arrive at random intervals where the gaps follow an exponential distribution. With $\lambda=0.5/s$ over 100s, this generates ~50 workflows. The exact count varies due to randomness.

Phase 5: Run Event Loop

```
scheduler.run(orchestrator::process);
```

One line, but this IS the entire simulation. The scheduler drains its priority queue, calling `orchestrator.process()` for each

event. Processing events schedules new events. The loop runs until no more events exist.

Phase 6: Export & Summary

Writes trajectory CSV and prints summary statistics: completed workflows, success rate, average/max latency, total cost, total steps.

16. Entry Point

File: SimulatorApplication.java

Spring Boot application with two modes:

- Default (no args): runs v1 Kubernetes microservice simulation.
- With "agentic" argument: runs v2 agentic simulation.

The agentic path simply creates an AgenticSimulationRunner and calls run().

```
if (args contains "agentic") {  
    new AgenticSimulationRunner().run("agentic_config.json");  
} else {  
    // v1 Kubernetes simulation  
}
```

17. End-to-End Flow Example

Let's trace a single workflow through the entire system with the default config: 1 agent (agent-1) in CLOUD, 1 LLM (gpt-4o) in CLOUD, 1 tool (search-tool) in CLOUD, intra-cloud latency = 1ms.

t=0.000ms: Workflow Submitted

Poisson traffic generator creates workflow 'wf-0001' with 50 prompt tokens. Orchestrator creates Message(USER -> agent-1, 50 tokens) and schedules AgentReceive at t=0.

t=0.000ms: Agent Receives Message (Step 1)

AgentService.onReceive() checks concurrency: 0 active < 10 max. Passes.
processReceive():

- Creates WorkflowContext(startedAt=0)
- stepIndex = 1, accumulatedInputTokens = 50
- Budget check: 1 <= 10 steps, 50 <= 10000 tokens. OK.
- Looks up LLM profile: gpt-4o in CLOUD
- Network: CLOUD->CLOUD = 1ms out + 1ms back = 2ms
- Samples outputTokens ~ Gaussian(150, 50) = e.g., 142
- Samples decision: roll=0.35 < 0.6 (call_tool weight) -> CALL_TOOL('search-tool')
- Samples TTFT ~ Gamma(4, 400) = e.g., 380ms
- Samples TPOT ~ Gamma(10, 25) = e.g., 24ms
- Inference = 380 + 142 * 24 = 3788ms
- Cost = 50 * \$2.5e-6 + 142 * \$10e-6 = \$0.001545
- Schedules LlmComplete at t = 0 + 1 + 3788 + 1 = 3790ms

t=3790ms: LLM Completes

Decision was CALL_TOOL('search-tool'). Agent dispatches tool call:

- Tool profile: search-tool in CLOUD
- Network: 1ms + 1ms = 2ms
- Tool latency ~ Gaussian(200, 50) = e.g., 185ms
- Error check: random() = 0.72 >= 0.0 (errorRate) -> no error
- Response tokens = 300
- Schedules ToolComplete at t = 3790 + 1 + 185 + 1 = 3977ms

t=3977ms: Tool Completes

Tool succeeded. Creates follow-up Message(search-tool -> agent-1, 300 tokens). Calls processReceive() directly (reuses concurrency slot).

t=3977ms: Agent Receives Tool Result (Step 2)

The tool result arrives back at the agent. processReceive() is called again, which triggers a NEW LLM call so the LLM

can evaluate the tool result and decide what to do next.

processReceive():

- stepIndex = 2, accumulatedInputTokens = 50 + 300 = 350
- Budget check: $2 \leq 10$, $350 \leq 10000$. OK.

t=3977ms: LLM Call #2 Dispatched

The agent sends the tool result (350 accumulated tokens) to the LLM for evaluation. This is a full LLM inference call with its own latency and cost:

- Samples outputTokens ~ Gaussian(150, 50) = e.g., 168
- Samples decision: roll=0.71 $\geq$ 0.6 (generate_text weight) -> GENERATE_TEXT
- Samples TTFT ~ Gamma(4, 400) = e.g., 390ms
- Samples TPOT ~ Gamma(10, 25) = e.g., 24ms
- Inference = $390 + 168 * 24 = 4422$ ms
- Cost = $350 * \$2.5e-6 + 168 * \$10e-6 = \$0.002555$
- Schedules LlmComplete at $t = 3977 + 1 + 4422 + 1 = 8401$ ms

Note: This is the same LLM call mechanism as Step 1. Every time the agent receives a message (whether from a user or a tool result), it makes an LLM call to decide the next action. The LLM uses the same probability weights (60% call_tool, 40% generate_text) each time.

t=8401ms: LLM Completes -> SUCCESS

The LLM's decision is GENERATE_TEXT, meaning it determined the tool result was sufficient to answer the user's question. Workflow terminates:

- Total latency: $8401 - 0 = 8401$ ms (8.4 seconds)
- Total cost: $\$0.001545 + \$0.002555 = \$0.0041$
- Total steps: 2
- Reason: SUCCESS

If the decision had been CALL_TOOL instead, the agent would have dispatched another tool call and the loop would continue (Step 3, Step 4, etc.) until the LLM finally decides GENERATE_TEXT or the budget is exhausted.

Concurrency slot freed. Next queued workflow (if any) is dequeued.

t=8401ms: Orchestrator Logs COMPLETE

TrajectoryCollector records the final row with all metrics. This workflow generated ~10 trace events total:

1. SUBMIT - workflow enters the system
2. AGENT_RECEIVE - agent gets user message
3. LLM_DISPATCH - LLM call #1 sent (agent -> LLM)
4. LLM_COMPLETE - LLM call #1 returns: CALL_TOOL
5. TOOL_DISPATCH - tool call sent (agent -> tool)
6. TOOL_RETURN - tool result comes back

- 7. AGENT_RECEIVE - agent receives tool result
- 8. LLM_DISPATCH - LLM call #2 sent (agent -> LLM)
- 9. LLM_COMPLETE - LLM call #2 returns: GENERATE_TEXT
- 10. COMPLETE - workflow finished

Notice how there are TWO LLM calls: one for the initial decision (step 1) and one after the tool result comes back (step 2). Every agent step involves an LLM call.

18. Key Design Decisions

Why Discrete Event Simulation?

DES is the gold standard for modeling queuing systems. It is:

- Exact: no approximation errors from time-stepping
- Fast: jumps between events, doesn't tick through empty time
- Deterministic: same seed = same results = reproducible experiments
- Scalable: can simulate thousands of workflows in seconds

Why Gamma for LLM Latency?

Real LLM inference latency measurements show:

1. Always positive (no negative latency)
2. Right-skewed (most calls are fast, some are slow)
3. The tail varies by model/provider

Gamma distribution matches all three properties. The shape parameter controls the tail: low shape = heavy tail (more outliers), high shape = Gaussian-like.

Why Token-Based (Not Byte-Based)?

In agentic systems, the fundamental unit is the TOKEN:

- LLM latency scales with output tokens
- Cost is per-token
- Context window limits are in tokens
- Tool results are measured in tokens

v1 used bytes (for network bandwidth simulation). v2 uses tokens throughout.

Why Sealed Events?

Java sealed interfaces guarantee exhaustive matching. If you add a new event type, the compiler forces you to handle it everywhere. This prevents bugs where a new event type silently falls through.

Why Pluggable Decision Policy?

The LLMEngine takes a BiFunction for decisions, not a hardcoded strategy. This lets you:

- Use weighted random for statistical experiments
- Use scripted sequences for debugging (e.g., 'always call tool twice then finish')
- Swap in ML-based policies in the future
- Test edge cases (e.g., 'always fail')

Concurrency = Queuing Theory

The maxConcurrency + queue model is a direct M/M/c queue:

- M: Poisson arrivals (Markovian)
- M: Exponential service times (approximately)
- c: number of servers (maxConcurrency)

This lets you study queue buildup, tail latency under load, and optimal concurrency settings for different arrival rates.